

David Barker

**Functional Reactive
Programming for Data Binding
in C#**

Computer Science Tripos Part II

Jesus College

March 1, 2013

Proforma

Name: **David Barker**
College: **Jesus College**
Project Title: **Functional Reactive Programming for
Data Binding in C#**
Examination: **Computer Science Tripos Part II**
Word Count: **TBC**
Project Originator: **Tomas Petricek**
Supervisor: **Tomas Petricek**

Original Aims of the Project

Aims of the project will go here.

Work Completed

Work completed will go here.

Special Difficulties

None

Declaration

I, David Barker of Jesus College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed [signature]

Date [date]

Contents

1	Introduction	1
1.1	Data binding	1
1.2	Data binding in .NET	2
1.3	Project inspirations	2
2	Preparation	5
3	Implementation	7
3.1	Overview	7
3.2	Arrows	7
3.2.1	Overview	7
3.2.2	Simple arrows	7
3.2.3	Invertible arrows	8
3.2.4	List arrows	8
3.2.5	Choice arrows	8
3.2.6	Further utility arrows	8
3.2.7	Feedback in arrows	8
3.3	Data binding	8
3.3.1	Overall architecture	8
3.3.2	Creating bindable sources and destinations	9
3.3.3	Creating bindings	10
	Syntax	10
	Cycle and conflict detection	10
	Two-way binding	10
	Many-to-many bindings	10
	Problems encountered	11
	Type safety	11
	Binding to lists	11
3.3.4	Integration work with WPF	11

4	Evaluation	13
4.1	Correctness of arrow implementations	13
4.1.1	Automated testing	13
	Simple arrows	13
	Invertible arrows	13
4.1.2	Correctness proof by decomposing into lambda calculus . .	13
4.2	Syntax evaluation	14
4.2.1	Arrow syntax	14
	Comparison with Haskell	14
4.2.2	Binding syntax	14
	Username two-way binding	14
	List binding from a mock database	14
	Some other demo	14
4.3	Performance testing	14
4.3.1	Arrow performance	14
	Measuring technique	14
	Simple function results	14
	List function results	14
	Overhead due to arrow chaining	15
4.3.2	Binding performance	16
5	Conclusion	17

List of Figures

3.1	Marshalling a list of BindPoint sources to pass to an arrow and unmarshalling the result	11
4.1	Performance of arrows, Funcs and normal functions in implementing simple functionality	15
4.2	Performance of arrows, Linq queries and normal (loop-based) functions in implementing simple list functionality	15

Chapter 1

Introduction

It has become a firmly established practice when writing applications to separate the code managing the user interface from the business logic powering it. This is a key part of the principle of separation of concerns, and means that the interface can be safely changed without needing to modify the backend code (and vice versa, assuming the backend provides a consistent interface). Beginning with the traditional MVC architecture, this has led to the development of a family of system architectures such as MVPM and MVVM which enforce this principle.

1.1 Data binding

Data binding presents a mechanism for bridging the gap between the separated layers by allowing the developer to specify that some value in the user interface code should be bound to a property of the model. This usually simply means that whenever the value in the model changes, the data binding framework will ensure that the value in the user interface is also updated to the same value. However, bindings can often be more complex: for instance, there could be a text box in the user interface which reflects backend data but can also be used to change it (two-way binding). Alternatively, there might be a user interface component whose value is determined by some function of a variable in the model - consider a list view which displays a filtered and sorted version of a list stored in the model. There could even be values which depend on multiple sources, or more complex many-to-many bindings.

Many current languages and frameworks provide data binding features with varying levels of complexity. Java has a variety of extension libraries [examples] which allow the programmer to use data binding with Swing. Javascript too has a variety of possibilities, with a prominent example being the backbone.js MVC

framework. Unfortunately, many of these are quite limited and difficult to use - backbone.js, for example, only provides one-way binding and requires a lot of boilerplate code from the programmer.

1.2 Data binding in .NET

Microsofts .NET framework offers a particularly powerful example of data binding through Windows Presentation Foundation (WPF). Based on the MVVM architecture, it has many features to allow things like two-way binding, binding through functions and bindings based on list operations. One of its key advantages is that the user interface can be defined entirely in XAML ¹ with bindings being specified through XAML parameters. The view logic is then specified in the ViewModel which in turn communicates with the model. This means user interface designers can work purely in XAML without concern for the logic or binding mechanisms in place behind the scenes.

However, WPF suffers from a similar problem to many other data binding frameworks: advanced data bindings can be very complex and difficult to manage, and setting up bindings in the first place requires quite a lot of boilerplate code in the model and view. Furthermore, binding through functions requires special value converter classes to be written. This is essentially an application of the Template pattern, and the value converters are not type safe - they take objects as input and return objects as output (and bindings with multiple inputs will simply take arrays of objects with no guarantee that the right number of values has been passed). [Maybe more disadvantages?] Clearly a more simple and general binding framework would definitely make application development simpler.

1.3 Project inspirations

Many useful ideas for data binding come from the area of functional programming. Functional reactive programming (FRP), for example, is a paradigm which nicely models the concept of data binding in a general way. Another useful concept is that of the arrows implemented in Haskell. An arrow from type A to type B essentially represents a process taking an input of type A and returning something of type B, and these can be arbitrarily combined in interesting ways to build up complex functions. More detail on both will be given in the next chapter.

¹Extensible Application Markup Language, an XML-based markup language for defining user interfaces and simple behaviour

The FRP paradigm was the inspiration for the project, which sought to implement a general-purpose data binding framework in C# using concepts derived from functional programming. I successfully implemented a framework which provides data binding in both directions, through arbitrary functions with arbitrary numbers of inputs and outputs, all with a simple syntax based on lambda expressions. Minimal boilerplate code is required to set up bindable properties - the containing class should extend Bindable, and from there marking arbitrary member variables and properties with the [Bindable] tag will make it possible to create bindings to and from them with a single function call. I also implemented a large variety of arrow types and arrow operators along with a number of utility functions to make creating and combining them simple.

Chapter 2

Preparation

This chapter is empty still!

Chapter 3

Implementation

3.1 Overview

Give a general overview

3.2 Arrows

3.2.1 Overview

Overview of how arrows worked

- Type inference chosen instead of reflection
- Techniques for allowing type inference - static constructor methods, extension methods etc.
- Arrow intercompatibility?
- Multiple inputs/outputs and the decision to use C# Tuple class

3.2.2 Simple arrows

Summarise functionality, combinators and syntactic efforts made here, mention things like the parallelism of the And operator as well.

Challenges encountered

A few - mention type inference challenges, overcoming type system whilst maintaining type safety and difficulties with the And operator (also alternatives considered for this)

3.2.3 Invertible arrows

Discuss the alternative ways of implementing this - a set of preset arrows designed to form a Turing-complete basis, forcing the user to specify an arrow for each direction etc. Final decision to simply construct using a function for each direction then allow arbitrary combination, as inspired by the paper on invertible arrows.

3.2.4 List arrows

- Wrapping Linq
- SQL-style syntax - helped by various extension methods
- Functional-style operations - foldl and foldr

3.2.5 Choice arrows

- Not deeply implemented or tested but could be handy for some things, most likely convenient exception handling
- This wouldn't be much good for bindings though

3.2.6 Further utility arrows

Identity arrows Exist for normal arrows and invertible ones; mention the reason for the class - has to be type parameterised

Tuple reassociation arrows Useful for stuff like arrow law tests, explain here...

Others Any more to mention?

3.2.7 Feedback in arrows

Discuss whether this would be useful or not, reasons for it not being there (lack of real-world use cases?) and how one might implement it.

3.3 Data binding

3.3.1 Overall architecture

Use diagrams and stuff to explain it.

3.3.2 Creating bindable sources and destinations

One of the main problems with WPF data binding is the complexity of making sources 'bindable'. This usually requires that the programmer manually implement the `INotifyPropertyChanged` interface, creating appropriate events and overriding the set methods of all their properties such that they throw the right events. This often becomes a case of copying and pasting the code used for other data sources as it is almost always identical and includes boilerplate tasks like ensuring a variable's new value is different from its old one before throwing the event, and checking that the event isn't null.

For this project I decided to abstract these details away into a base class, called `Bindable`. This provides events for binding and a set of methods used by binding classes to manage data binding, many of which will be discussed later in this section. The most important are those which allow getting and setting of arbitrary variables by name using reflection, as these are the main interface used by the bindings manager. These have been designed to be dynamically type safe (that is, throwing exceptions at run time where type errors are encountered), to ensure the properties being accessed exist and to abstract away the differences between properties and simple public member variables¹.

The result of this is that instead of writing all the boilerplate code one would usually write, the programmer need only make their sources and destinations extend `Bindable` and all the usual things will be handled elsewhere leading to considerably less code clutter. A representative example of the final syntax is given in Listing 3.1.

Another addition that can be seen in the listing is the `[Bindable]` attribute which has been used above the `value` property. This uses a `PostSharp` aspect, explained in the next section, to intercept the setter for throwing events and so relieve the programmer of having to do this themselves.

Listing 3.1: Creating a bindable source class

```
public class DataSource : Bindable
{
    [Bindable]
    public int value { get; set; }
}
```

¹WPF does not allow binding to member variables, but for this project I decided there was no obvious need to distinguish between the two

PostSharp

PostSharp is a C# extension which provides aspect-oriented programming (AOP) functionality. Essentially, AOP is a means of separating cross-cutting concerns from the code they affect, allowing 'aspects' to be written providing this common functionality. Markers can then be placed in the appropriate points, and the aspects will generate and inject their code into these locations at compile time.

In this case, the `Bindable` aspect will intercept the setter of any property or member variable it is placed above and insert code to check whether the value has changed and throw appropriate events if it has.

3.3.3 Creating bindings

Syntax

Show what it looks like.

Cycle and conflict detection

Techniques used and all that.

Two-way binding

Problems encountered - need for variable locking system to avoid infinite loop

Many-to-many bindings

Many-to-many bindings presented an additional challenge for the binding system. It was decided that the best approach in terms of syntax would be to have the user pass in a list of source bind points, an arrow and a list of destination bind points, so as to mirror the syntax used in simple bindings. However, this simple approach is out of step with the way parameters are passed to arrows: by construction, all arrows taking multiple inputs will take them in the form of a binary-tree-structured Tuple². As there didn't appear to be any C# syntax allowing the user to pass in binding sources in a similar fashion (such that the structure could be understood by the binding), it was decided that an argument marshaller/unmarshaller would be needed to handle the translation between lists of bind points and tree-structured tuples.

²This can be verified by looking at the implementations of the operators; for instance, `And` converts two arrows (which may be on tuples) into an arrow on two-tuples containing the types of the original arrows, thus building up a binary tree structure

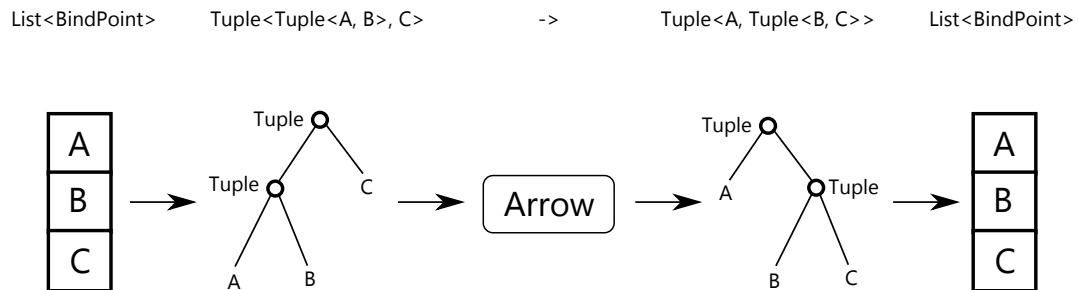


Figure 3.1: Marshalling a list of BindPoint sources to pass to an arrow and unmarshalling the result

By using reflection on the type of the arrow, I was able to convert a list of bind points to a binary-tree-structured Tuple. The code essentially does a recursive depth-first search through the input type of the arrow (obtained via reflection), simultaneously constructing a Tuple of the same structure and inserting values from the bind points wherever a 'leaf' is encountered in the type being searched. At the other end, a similar process is used to extract the values from the returned tuple and put them into a list which is then used to assign to the destination bind points. The process is illustrated in Figure 3.1, where a set of inputs of types A, B and C are being passed into an arrow from type `Tuple<Tuple<A, B>, C>` to type `Tuple<A, Tuple<B, C>>`.

Problems encountered

Type safety All sort of problems here; difficulties of type checking arrows should also be mentioned.

Binding to lists They don't throw events! Binding to lists needs the list to be assigned to for it to update. Possible solutions like `ObservableEnumerable` (or whatever it was called), potentially more boilerplate, future work n shit.

3.3.4 Integration work with WPF

Mention difficulties with the value converter and inability to properly integrate it into XAML and stuff.

Chapter 4

Evaluation

4.1 Correctness of arrow implementations

Mention here that the main device for determining their correctness has been the arrow laws, and possibly list them with explanations? Also mention that 'special' arrows like ListArrows are correct because they are based on simple arrows and invertible arrows. Also, maybe say that ArrowChoice was ignored here as it's slightly extraneous and its correctness probably follows from the correctness of simple arrows anyway.

4.1.1 Automated testing

Simple arrows

Give tests and results for arrows.

Invertible arrows

Same for invertible arrows.

4.1.2 Correctness proof by decomposing into lambda calculus

Explain the technique used to translate into lambda calculus and the likely correctness of this, and give a few sample proofs for some non-trivial arrow laws.

4.2 Syntax evaluation

4.2.1 Arrow syntax

Arrow syntaxy stuff.

Comparison with Haskell

Might not be great, but why not?

4.2.2 Binding syntax

Pretty good. Say something about the demo apps and stuff:

Username two-way binding

Todo

List binding from a mock database

Todo

Some other demo

Todo

4.3 Performance testing

4.3.1 Arrow performance

Tests indicating how arrows perform vs. functions and funcs.

Measuring technique

Explain how the results were obtained

Simple function results

The simple function results can be seen in the graph.

List function results

Results for list-based functions




Figure 4.1: Performance of arrows, Funcs and normal functions in implementing simple functionality




Figure 4.2: Performance of arrows, Linq queries and normal (loop-based) functions in implementing simple list functionality

Overhead due to arrow chaining

Explanation of the problem (which will have been highlighted in earlier results) along with test results for increasingly long chains of identity functions. Suggest reasons for this and potential solutions (or maybe the solutions should come in

later?)

4.3.2 Binding performance

Todo...

Chapter 5

Conclusion

Conclusion goes here!

